

CS FUNDAMENTALS

STUDY MODULE



Q&A Format with Practical Applications



Disclaimer

Everyone learns uniquely.

What matters is developing the problem solving ability to solve new problems.

This Doc will help you build strong foundational knowledge.



Topic-1

Networking Concepts

Question

Explain the OSI model with an emphasis on the functions of each layer.

Answer

The OSI model has seven layers:

- **Physical Layer:** Deals with the physical connection between devices.
- **Data Link Layer:** Responsible for node-to-node data transfer and error detection.
- **Network Layer:** Manages data routing, forwarding, and addressing.
- **Transport Layer:** Ensures end-to-end communication, flow control, and error recovery.
- **Session Layer:** Manages sessions in a network, controlling dialogues between computers.
- **Presentation Layer:** Translates data between the application and the network.
- **Application Layer:** Interfaces directly with user applications, enabling network services.

Question

What is subnetting, and how would you calculate the subnet mask for a given network?

Answer

Subnetting divides a network into smaller subnets to improve management and security. To calculate the subnet mask, determine the number of required subnets, convert this into binary, and subtract it from the default subnet mask. For example, for a network needing 8 subnets, subtract 3 bits from the host part of the IP address, resulting in a subnet mask of 255.255.255.224.

Question

Describe the process of DNS resolution from start to finish.

Answer

DNS resolution involves translating a domain name into an IP address. When a user types a URL, the request goes to a DNS resolver, which queries root servers, top-level domain (TLD) servers, and authoritative name servers. The resolver then retrieves the IP address and returns it to the user's browser, which connects to the web server using this IP.

Question

What is Quality of Service (QoS) and how is it implemented in a network?

Answer

QoS is a set of technologies used to manage and prioritize network traffic to ensure the performance of critical applications. It is implemented using traffic shaping, priority queuing, and bandwidth management techniques, typically on routers and switches.

Question

Explain how TCP's three-way handshake works and its importance in ensuring reliable communication.

Answer

The TCP three-way handshake is crucial for establishing a reliable connection between a client and server. It involves three steps:

- 1. SYN:** The client sends a synchronization packet to initiate the connection.
- 2. SYN-ACK:** The server acknowledges the SYN and responds with its own SYN.
- 3. ACK:** The client acknowledges the server's SYN-ACK, establishing the connection.

Practical Application:

Question

How would you design a corporate network for a company with multiple branches to ensure secure and efficient communication?

Answer

The design would include VPNs for secure communication between branches, VLANs to segregate traffic, and QoS to prioritize critical services. Implement redundancy with dual WAN links, use NAT for efficient IP address usage, and establish a centralized firewall to manage security across all branches.

Topic-2

Network Protocols and Security

Question

Explain the difference between TCP and UDP, and provide examples of applications that use each protocol.

Answer

TCP (Transmission Control Protocol) is connection-oriented, ensuring reliable data transmission with error checking and flow control (e.g., web browsing, email). UDP (User Datagram Protocol) is connectionless, offering faster data transmission without reliability checks (e.g., live streaming, online gaming).

Question

Describe how SSL/TLS encryption works in securing communications over the internet.

Answer

SSL/TLS uses asymmetric encryption to establish a secure session between client and server. During the handshake, the server provides a public key in its certificate, which the client uses to encrypt a session key. The session key is then used for symmetric encryption of the data transmitted between client and server.

Question

What is BGP, and why is it critical for the operation of the internet?

Answer

BGP (Border Gateway Protocol) is used to exchange routing information between different autonomous systems (AS) on the internet. It ensures that data can find the best path across multiple networks, making it critical for maintaining global internet connectivity and route optimization.

Question

Explain the concept of Network Address Translation (NAT) and its different types.

Answer

NAT translates private IP addresses to a public IP address, allowing multiple devices on a local network to share a single public IP for internet access. Types include Static NAT (one-to-one mapping), Dynamic NAT (many-to-one mapping), and PAT (Port Address Translation), which maps multiple private IPs to a single public IP using different ports.

Question

How does a VPN enhance network security, and what are the key protocols used in VPNs?

Answer

A VPN (Virtual Private Network) enhances security by encrypting data sent over public networks, creating a secure tunnel between the user's device and the VPN server. Key protocols include IPsec (Internet Protocol Security), SSL/TLS (for secure web-based VPNs), and PPTP (Point-to-Point Tunneling Protocol).

Question

What is a firewall, and how does it function in protecting a network?

Answer

A firewall is a security device that monitors and controls incoming and outgoing network traffic based on predetermined security rules. It functions by filtering traffic, blocking unauthorized access, and allowing only legitimate communications to pass through based on criteria such as IP addresses, protocols, and ports.

Practical Application:

Question

Design a secure network architecture for an e-commerce platform, considering high traffic and sensitive data.

Answer

The architecture would include a multi-tier design with a DMZ (Demilitarized Zone) for public-facing services like web servers, behind which are application and database servers protected by firewalls. Use SSL/TLS for encrypting customer data, implement intrusion detection/prevention systems (IDS/IPS), and deploy load balancers for traffic management. Regular backups and network segmentation would also be crucial for security and performance.

Topic-3

Advanced Networking Concepts

Question

What is MPLS and how does it differ from traditional IP routing?

Answer

MPLS (Multiprotocol Label Switching) is a technique that directs data from one network node to the next based on short path labels rather than long network addresses, avoiding complex lookups in a routing table. Unlike traditional IP routing, which forwards packets based on destination IP, MPLS can create predetermined, optimized paths through the network, improving performance.

Question

Describe the role of load balancers in high availability systems.

Answer

Load balancers distribute incoming network traffic across multiple servers, ensuring no single server becomes a bottleneck, thus enhancing system reliability and availability. They monitor server health, redirecting traffic away from failed servers to prevent downtime.

Question

Explain the concept of VLAN tagging and its importance in network segmentation.

Answer

VLAN tagging involves adding a tag to Ethernet frames to indicate the VLAN membership of a packet. This is important for network segmentation, allowing multiple VLANs to coexist on a single physical network infrastructure, improving security and traffic management by isolating different types of traffic.

Question

What is a distributed denial-of-service (DDoS) attack, and how can networks be protected against it?

Answer

A DDoS attack involves overwhelming a target network or server with a flood of traffic from multiple sources, leading to service disruption. Protection strategies include using firewalls and intrusion detection systems, deploying DDoS mitigation services, rate limiting, and ensuring redundancy in network architecture.

Question

How does Quality of Service (QoS) improve the performance of VoIP and video conferencing applications?

Answer

QoS improves performance by prioritizing traffic types, ensuring that real-time applications like VoIP and video conferencing receive the necessary bandwidth, low latency, and low jitter to function effectively. This is done by classifying and managing traffic according to predefined policies.

Question

What are the advantages and challenges of implementing SD-WAN in a corporate network?

Answer

SD-WAN (Software-Defined Wide Area Network) offers advantages like improved application performance, centralized management, and cost savings by using cheaper broadband connections. Challenges include complexity in initial setup, potential security risks if not configured correctly, and the need for reliable internet connections across all locations.

Practical Application:

Question

How would you implement SD-WAN in a multi-branch organization to optimize application performance and reduce costs?

Answer

Implement SD-WAN by deploying edge devices at each branch to route traffic intelligently across multiple WAN links based on real-time network conditions. Use centralized management software to define policies that prioritize critical applications and route traffic accordingly. Leverage broadband internet for non-critical traffic while reserving MPLS links for high-priority services, reducing overall costs.



DATABASE MANAGEMENT SYSTEMS (DBMS)

Topic-1

Intermediate to Advanced DBMS Concepts

Question

What are ACID properties, and why are they essential in a database system?

Answer

ACID properties—Atomicity, Consistency, Isolation, and Durability—are essential to ensure that database transactions are processed reliably. Atomicity ensures all operations within a transaction are completed; Consistency ensures that a transaction brings the database from one valid state to another; Isolation ensures that transactions do not interfere with each other; Durability ensures that the results of a transaction are permanently recorded.

Question

Explain the different types of database indexes and their use cases.

Answer

Database indexes can be clustered (where the data is stored in the order of the index) or non-clustered (where the index is separate from the data storage). Clustered indexes are used for primary key columns, while non-clustered indexes are used for quick searches in non-primary key columns. Full-text indexes are used for searching within text data, and bitmap indexes are efficient for columns with a limited range of values.

Question

Describe the process of normalization and denormalization in databases.

Answer

Normalization involves organizing a database to reduce redundancy and improve data integrity, typically through dividing tables and establishing relationships. Denormalization is the process of combining tables to improve read performance at the cost of introducing redundancy, often used in data warehousing where query performance is prioritized over update efficiency.

Question

What is a transaction isolation level, and why is it important?

Answer

Transaction isolation levels determine how transaction integrity is visible to other transactions and prevent issues like dirty reads, non-repeatable reads, and phantom reads. Common isolation levels include Read Uncommitted, Read Committed, Repeatable Read, and Serializable, each offering different balances between consistency and performance.

Question

Explain the concept of database sharding and its benefits.

Answer

Database sharding involves splitting a large database into smaller, more manageable pieces called shards, each stored on a separate server. This approach can improve performance and scalability by distributing the load across multiple servers, making it easier to manage and access data in large-scale systems.

Question

What are the differences between SQL and NoSQL databases, and when would you choose one over the other?

Answer

SQL databases are relational, using structured query language for complex queries and transactions where consistency is key.

NoSQL databases are non-relational, offering flexible data models and are preferred for scalability, high availability, and handling unstructured data.

Practical Application:

Question

How would you design a database schema for a social media application to efficiently store and retrieve user posts, comments, and likes?

Answer

Design a relational database schema with tables for Users, Posts, Comments, and Likes. Implement indexes on frequently searched columns like user ID and post ID to optimize queries. Use foreign keys to enforce relationships and consider denormalization for read-heavy operations. Alternatively, use a NoSQL database like MongoDB for flexible, scalable schema design for storing posts and comments.

Topic-2

SQL and Query Optimization

Question

Write an SQL query to find the nth highest salary in a company's employee table.

Answer

properties

```
SELECT DISTINCT Salary FROM Employee ORDER BY Salary  
DESC LIMIT 1 OFFSET n-1;
```

Replace n with the required rank (e.g., 2 for the second highest).

Question

Explain the concept of JOINs in SQL, and describe the different types of JOINs available.

Answer

A JOIN clause in SQL is used to combine rows from two or more tables based on a related column. Types of JOINs include:

- **INNER JOIN:** Returns only matching rows from both tables.
- **LEFT JOIN (or LEFT OUTER JOIN):** Returns all rows from the left table and matched rows from the right table.

- **RIGHT JOIN (or RIGHT OUTER JOIN):** Returns all rows from the right table and matched rows from the left table.
- **FULL OUTER JOIN:** Returns all rows when there is a match in either table.
- **CROSS JOIN:** Returns the Cartesian product of both tables.

Question

What is a stored procedure, and how can it improve database performance?

Answer

A stored procedure is a precompiled set of SQL statements stored in the database, executed as a single unit. It improves performance by reducing network traffic, enables reuse of complex logic, and enhances security by controlling data access.

Question

Describe how query optimization works in an RDBMS.

Answer

Query optimization in an RDBMS selects the most efficient execution plan for a SQL query. The query optimizer evaluates plans based on factors like indexes, join order, and data distribution. Techniques include using optimal indexes, cost-based optimization, and query rewriting.

Question

Explain the concept of a deadlock in a database and how it can be resolved.

Answer

A deadlock occurs when two or more transactions block each other by holding locks on resources that the other transactions need. It can be resolved by implementing deadlock detection algorithms, using timeouts, or ensuring proper transaction isolation levels to prevent circular wait conditions.

Question

What is a CTE (Common Table Expression), and how is it used in SQL queries?

Answer

A CTE is a temporary result set that can be referenced within a SQL SELECT, INSERT, UPDATE, or DELETE statement. It is defined using the WITH clause and is used to simplify complex queries by breaking them into more manageable parts, making them easier to read and maintain.

Practical Application:

Question

How would you optimize a query in a large database to retrieve user activity logs within a specific date range efficiently?

Answer

To optimize the query, ensure that the activity logs table is indexed on the date column to speed up range queries. Consider partitioning the table by date to reduce the amount of data scanned. Use a CTE to handle complex subqueries, and if the table is extremely large, consider archiving old data to reduce the table size.

Topic-3

Advanced DBMS Concepts

Question

What is the CAP theorem, and how does it apply to distributed databases?

Answer

The CAP theorem states that a distributed database can only achieve two out of three properties at the same time: Consistency (every read receives the latest write), Availability (every request receives a response), and Partition Tolerance (the system continues to operate despite network partitions). Depending on the application, a database system might prioritize one of these properties over the others.

Question

Describe how database replication works and its benefits.

Answer

Database replication involves copying and maintaining database objects, such as tables and transactions, across multiple database servers. Benefits include improved data availability, fault tolerance, and load balancing, as it allows for redundancy and distributed query processing.

Question

What is a materialized view, and when would you use it?

Answer

A materialized view is a database object that contains the results of a query. Unlike a regular view, which is a virtual table, a materialized view stores data physically, allowing for faster query performance, especially for complex aggregations and joins. It is used when performance is critical, and the underlying data does not change frequently.

Question

Explain the concept of eventual consistency in NoSQL databases.

Answer

Eventual consistency is a consistency model used in distributed databases where updates to the database are propagated asynchronously. The system guarantees that, given enough time, all replicas of the data will converge to the same state, but reads may return stale data until that convergence happens. It is commonly used in systems where availability and partition tolerance are prioritized over immediate consistency.

Question

What are database transactions, and how do transaction logs ensure data integrity?

Answer

Database transactions are units of work that are performed atomically, ensuring that either all operations within the transaction are completed successfully or none are applied at all. Transaction logs record all changes made during a transaction, enabling recovery and ensuring data integrity even in the event of a system crash.

Question

What is database sharding, and what are the challenges associated with it?

Answer

Sharding is the process of splitting a database into smaller, distributed pieces called shards, each stored on a separate server. The challenges include managing cross-shard queries, ensuring data consistency across shards, and handling shard rebalancing as data grows.

Practical Application:

Question

How would you design a distributed database system to handle high write and read volumes while ensuring data consistency?

Answer

Design the system using a combination of sharding for horizontal scaling, replication for fault tolerance and high availability, and a consensus algorithm like Raft or Paxos to ensure consistency across nodes. Implement eventual consistency if the application can tolerate it, or use strong consistency models for critical operations. Use a distributed database like Cassandra or MongoDB that supports these features.



Topic-1

Intermediate to Advanced OS Concepts

Question

Explain the difference between a process and a thread.

Answer

A process is an independent program in execution with its own memory space, while a thread is a smaller execution unit within a process sharing the same memory space. Threads can communicate easily but require synchronization to avoid conflicts.

Question

What is virtual memory, and how does it work?

Answer

Virtual memory is a technique that gives applications the illusion of contiguous memory, even if it's fragmented across physical locations. The OS uses paging and swapping to manage virtual memory, enabling larger applications to run despite limited physical memory.

Question

Describe the concept of deadlock and the conditions necessary for a deadlock to occur.

Answer

A deadlock is a situation where a set of processes are blocked because each process is holding a resource and waiting for another resource held by another process. The conditions necessary for deadlock are Mutual Exclusion, Hold and Wait, No Preemption, and Circular Wait.

Question

What are page replacement algorithms, and how do they impact system performance?

Answer

Page replacement algorithms decide which memory pages to swap out when a new page needs to be loaded into memory. Common algorithms include FIFO (First In, First Out), LRU (Least Recently Used), and LFU (Least Frequently Used). The choice of algorithm impacts system performance by affecting cache hit rates and the frequency of page faults.

Question

How do modern operating systems manage multiple processes and ensure fair CPU allocation?

Answer

Modern operating systems use scheduling algorithms like Round Robin, Priority Scheduling, and Multilevel Queue Scheduling to manage multiple processes. These algorithms decide the order in which processes are executed by the CPU, ensuring fair allocation of CPU time based on process priority, time quantum, and other factors.

Question

What is the purpose of system calls in an operating system? Provide examples.

Answer

System calls are the interface between user programs and the operating system. They provide a way for programs to request services from the OS, such as file operations, process control, and network communication. Examples include `fork()` for creating a new process, `read()` for reading data from a file, and `socket()` for network communication.

Practical Application:

Question

Design an OS-level solution for handling high levels of concurrency in a web server application.

Answer

Implement a multithreaded server using thread pools to manage concurrent connections efficiently. Use non-blocking I/O operations and epoll or select system calls for managing large numbers of simultaneous connections. Ensure proper synchronization mechanisms like mutexes or semaphores are in place to avoid race conditions. Optimize CPU scheduling to prioritize critical server tasks.

Topic-2

Memory Management

Question

What is virtual memory, and why is it important in modern operating systems?

Answer

Virtual memory lets an OS use more memory than available by swapping RAM to disk. It simplifies memory management, isolates processes, and enhances security by preventing unauthorized memory access.

Question

Explain the concept of paging and how it differs from segmentation.

Answer

Paging divides memory into fixed-size pages, mapping logical to physical addresses, simplifying management but causing internal fragmentation. Segmentation divides memory into variable-sized segments based on logical divisions, which is more intuitive but may lead to external fragmentation.

Question

Describe the purpose and implementation of a page table in memory management.

Answer

A page table is a data structure used in virtual memory systems to map logical addresses to physical addresses. It keeps track of where each page of a process is located in physical memory or on disk. The OS uses the page table during address translation, and it can be hierarchical, inverted, or hashed to optimize performance and reduce memory overhead.

Question

What are the benefits and drawbacks of using a multilevel page table?

Answer

A multilevel page table reduces the memory overhead of storing large page tables by breaking them into smaller, hierarchical tables. This allows for more efficient use of memory. However, the drawback is increased complexity in address translation, as multiple memory accesses are required to resolve a single logical address, potentially impacting performance.

Question

Explain the difference between demand paging and pre-paging in operating systems.

Answer

Demand paging loads pages into memory only when they are needed, reducing initial memory usage but potentially leading to higher page fault rates. Pre-paging attempts to predict and load pages into memory before they are needed, reducing page faults but potentially wasting memory on pages that are not used.

Question

What is a TLB (Translation Lookaside Buffer), and how does it enhance memory management?

Answer

A TLB is a cache used by the CPU to reduce the time taken to access the page table during address translation. It stores the most recent translations of logical addresses to physical addresses, allowing for faster memory access and reducing the number of page table lookups, which enhances overall system performance.

Practical Application:

Question

Design a memory management strategy for a real-time operating system where low latency is critical.

Answer

Use a combination of fixed-size memory allocation and real-time paging with a locked memory approach to ensure low latency. Implement a TLB with minimal misses by keeping critical pages resident in memory. Use a priority-based paging strategy where high-priority tasks are guaranteed memory access. Consider the use of pre-allocated memory pools for critical tasks to avoid dynamic memory allocation overhead.

Topic-3

Process Synchronization and Deadlocks

Question

What are semaphores, and how are they used in process synchronization?

Answer

Semaphores are synchronization tools that control access to shared resources in concurrent systems, preventing race conditions by limiting the number of processes accessing a critical section or resource at a time.

Question

Explain the concept of a monitor and how it differs from a semaphore.

Answer

A monitor is a higher-level synchronization construct that combines mutual exclusion with condition variables, automatically handling locking and unlocking for safer, structured synchronization.

Question

Describe the readers-writers problem and how it can be solved.

Answer

The readers-writers problem deals with scenarios where multiple threads or processes need to read and write shared data. The challenge is to ensure that no writer is writing while another process is reading or writing. Solutions include using semaphores to give priority to either readers or writers, depending on the application's needs, or employing a read-write lock that allows multiple readers but exclusive access for writers.

Question

What are the conditions that lead to a deadlock, and how can deadlocks be prevented?

Answer

The conditions for deadlock are Mutual Exclusion, Hold and Wait, No Preemption, and Circular Wait. Deadlocks can be prevented by ensuring that at least one of these conditions is not met.

Techniques include acquiring all required resources at once (eliminating Hold and Wait), using a timeout to preempt resource holding, and ordering resource allocation to prevent circular wait.

Question

Explain how the banker's algorithm works for deadlock avoidance.

Answer

The Banker's algorithm is a deadlock avoidance strategy that checks if resource allocation will leave the system in a safe state. Before granting a resource request, the algorithm simulates the allocation and checks if all processes can still complete with the available resources. If they can, the request is granted; otherwise, it is denied to avoid a potential deadlock.

Question

What is the difference between a livelock and a deadlock?

Answer

In a deadlock, processes are stuck waiting for each other indefinitely, while in a livelock, processes continuously change their state in response to each other but do not make progress, essentially "livelocking" each other out of completing their tasks. Livelocks can be resolved by implementing a backoff strategy where processes pause or retry operations with delays.

Practical Application:

Question

Design a synchronization mechanism for a multi-threaded application that processes transactions where consistency is critical.

Answer

Implement a monitor with condition variables to manage access to the shared transaction data, ensuring that only one transaction is processed at a time. Use semaphores to manage thread access to the monitor, with priority given to threads processing high-priority transactions. Employ a deadlock detection algorithm to monitor and resolve potential deadlocks by rolling back and retrying transactions in a safe order.



FINAL SECTION: MAANG INTERVIEW QUESTIONS

15 Most Commonly Asked Questions:

Question

Explain the concept of CAP theorem in distributed systems.

Answer

The CAP theorem states that in a distributed system, it is impossible to simultaneously guarantee Consistency, Availability, and Partition Tolerance. Systems must choose between prioritizing consistency or availability when a network partition occurs.

Question

How does a database index work? What are the types of indexes?

Answer

A database index works by creating a data structure that improves the speed of data retrieval. Types include clustered, non-clustered, full-text, and bitmap indexes, each optimized for different types of queries and data.

Question

Explain the difference between multithreading and multiprocessing.

Answer

Multithreading involves running multiple threads within the same process, sharing memory and resources, while multiprocessing involves running multiple processes, each with its own memory space, enabling parallel execution.

Question

Describe the process of DNS resolution.

Answer

DNS resolution converts domain names into IP addresses through a hierarchical process involving DNS resolvers, root servers, TLD servers, and authoritative name servers.

Question

What is the difference between Docker and a Virtual Machine?

Answer

Docker containers share the host OS kernel and are lightweight, while Virtual Machines (VMs) run full OS instances on virtualized hardware, making them more resource-intensive.

Question

How would you design a URL shortening service?

Answer

Design a distributed database for mappings, an API for short URLs, and a load balancer for traffic. Use hashing for unique URLs and implement caching for better performance.

Question

Explain the difference between HTTP and HTTPS.

Answer

HTTP is unencrypted, while HTTPS uses SSL/TLS encryption to secure data transmission between the client and server, protecting against eavesdropping and man-in-the-middle attacks.

Question

What are the different types of locks in a database? Explain optimistic and pessimistic locking.

Answer

Locks include shared, exclusive, and update locks. Optimistic locking checks for conflicts only at commit time, while pessimistic locking prevents conflicts by locking resources early in the transaction.

Question

How does load balancing work, and what are the strategies?

Answer

Load balancing distributes traffic across servers using strategies like round-robin, least connections, IP hash, and weighted balancing to optimize resource use and prevent overload.

Question

Describe the OSI model and the functions of each layer.

Answer

The OSI model consists of seven layers: Physical, Data Link, Network, Transport, Session, Presentation, and Application, each responsible for specific network functions, from physical data transmission to application-level interactions.

Question

What is RAID? Explain different RAID levels.

Answer

RAID (Redundant Array of Independent Disks) combines multiple drives for redundancy and performance, with common levels including RAID 0 (striping), RAID 1 (mirroring), RAID 5 (striping with parity), and RAID 10 (combining RAID 1 and 0).

Question

Explain the concept of consistency, availability, and partition tolerance in distributed systems.

Answer

Consistency ensures all nodes see the same data at the same time, availability guarantees that every request receives a response, and partition tolerance allows the system to continue operating despite network partitions.

Question

How do garbage collection and memory management work in Java?

Answer

Garbage collection in Java automatically reclaims memory by identifying and disposing of objects no longer in use. The JVM uses different algorithms (e.g., mark-and-sweep, generational garbage collection) to manage memory efficiently.

Question

What is the difference between a process and a thread? When would you use one over the other?

Answer

A process is an independent program with its own memory space, while a thread is a smaller unit within a process that shares memory with other threads. Use processes for isolation and security, and threads for tasks requiring shared resources and faster communication.

Question

Design a scalable system for a messaging application like WhatsApp. (System Design Question)

Answer

The system should include a distributed database for storing messages, a microservices architecture for handling different features (e.g., messaging, media sharing), load balancers, and messaging queues for asynchronous processing. Use encryption for data security and implement redundancy and failover mechanisms for high availability.



WHY BOSSCODER?

 **1000+ Alumni** placed at Top Product-based companies.

 More than **136% hike** for every 2 out of 3 Working Professional.

 Average Package of **24LPA**.

The syllabus is most up-to-date and the list of problems provided covers all important topics.

Lavanya
 Meta



Course is very well structured and streamlined to crack any MAANG company

Rahul
 Google



[EXPLORE MORE](#)